

MERGE_{NET}: DIFFERENTIABLE SPATIAL ROUTING WITH A PHYSICAL TOKEN BOTTLENECK

Anonymous authors

Paper under double-blind review

ABSTRACT

Vision transformers preserve fine image detail but pay to process every patch through global blocks. Learned token merging must decide where information flows while presenting a shorter sequence to later blocks. DTEM learns merge decisions through a continuous relaxation, but uses separate soft and hard training paths to obtain physical compression. We introduce MergeNet, a local-to-latent transformer that combines differentiable spatial mass transport with an exact-budget gather in one forward pass. Selected carriers read full-grid features through cross-attention and retain their transported mass in latent attention. A rank-one query/key augmentation makes this mass bias compatible with FlashAttention. On ImageNet-1K with DeiT-S/8 at 392 final patches, MergeNet reaches 81.360% top-1, compared with 81.932% for post-training ToMe and 82.246% for dense DeiT-S/8. In 150-epoch routing comparisons, radius-3 exceeds global and flattened-window routing by 1.670 and 0.598 percentage points; three-seed CIFAR-100 controls also favor spatial routing. The matched DTEM-p8 adaptation reaches [VALUE]%. Under a common trained-checkpoint protocol, the accuracy–efficiency comparison is [comparison summary].

1 INTRODUCTION

Vision transformers use patch tokens to represent image detail and self-attention to integrate information across the image. Processing a fine patch grid through every block, however, makes this global reasoning expensive. Token merging reduces the sequence by combining patches rather than simply removing them (Bolya et al., 2023). A learned merger thus faces two linked requirements: learn where information should flow, and physically reduce the sequence seen by downstream blocks during training.

Differentiable merging does not automatically solve this execution problem. DTEM learns dedicated merge embeddings through continuous grouping and merging (Lee & Hong, 2024). Its relaxed operators retain all token slots; for end-to-end training, DTEM alternates soft updates of the merge embeddings with hard, token-reducing updates of the backbone. Thus DTEM uses both continuous policy learning and physical compression during training, but in separate update paths. We ask whether one forward pass can propagate gradients through soft routing while downstream blocks process an exact, physically shorter sequence.

We study this question with MergeNet, a local-to-latent vision transformer. Six local blocks encode the full patch grid, after which soft mass transport learns how patch information moves between slots. A hard gather then selects K carriers from N patches, so the subsequent latent blocks receive a physically shorter tensor during both training and inference. A residual cross-attention read lets those carriers access the full-grid local features before global latent reasoning. The gather indices are discrete; a straight-through weight connects the selected features to continuous routing scores without claiming to differentiate the indices themselves.

The routing support follows the original two-dimensional patch grid. A donor can send mass only to eligible receivers within radius R , while learned similarities determine the weights and the resulting mass determines which carriers survive the gather. We compare this original-grid routing with global, flattened-window, and degree-matched alternatives.

Carrier mass also has a defined role after selection. Following proportional attention in ToMe (Bolya et al., 2023), MergeNet adds each carrier’s $\log m$ to its attention key logits. This per-key bias is rank one across queries; the released operator reproduces it by adding one query/key dimension and using the original softmax scale in an unmodified FlashAttention call (Dao et al., 2022). We use an existing threshold-based smooth top- k relaxation (Su, 2024) for the continuous selection scores. Our contribution is the way these components connect soft spatial transport to an exact-budget latent path, together with the mass-bias implementation.

The experiments characterize the spatial-routing variant, with a clear limit on the overall performance claim. In a single-seed, 150-epoch ImageNet-1K study, radius-3 routing scores 1.670 percentage points above the global variant and 0.598 points above the flattened-window variant. A three-seed CIFAR-100 study also favors radius 3 over global, flattened, and degree-matched routing. At a matched final budget of 392 patches, the trained 300-epoch MergeNet checkpoint reaches 81.360% ImageNet top-1, below post-training ToMe’s 81.932%; the uncompressed DeiT-S/8 reaches 82.246%. The matched DTEM-p8 adaptation reaches [VALUE]% top-1. Together, the routing and cross-method results establish the behavior of the physical bottleneck; the common-protocol accuracy–latency comparison is [comparison summary].

Our contributions are:

- a local-to-latent architecture that combines differentiable spatial mass transport with an exact-budget gather before latent processing;
- comparisons of original-grid routing with global, flattened-window, and degree-matched alternatives, including a three-seed CIFAR-100 study; and
- an exact query/key reparameterization for mass-aware attention with an unmodified FlashAttention call, accompanied by separate reporting of accuracy, token budget, arithmetic cost, and measured runtime.

2 RELATED WORK

Token reduction. DynamicViT learns token retention (Rao et al., 2021), while A-ViT assigns adaptive depth through halting (Yin et al., 2022). ToMe uses bipartite soft matching and can compress pretrained or jointly trained transformers (Bolya et al., 2023); PiToMe (Tran et al., 2024), MCTF (Lee et al., 2024), and SAD-TM (Xie et al., 2026) modify its matching signal or merge schedule. DiffRate treats layerwise budgets as another learned variable (Chen et al., 2023). Equal final counts therefore do not imply equal layerwise computation.

Differentiable merging. DTEM is the closest differentiable predecessor. It learns lightweight embeddings dedicated to merge decisions, rather than choosing token pairs from the backbone features alone (Lee & Hong, 2024). A continuous relaxation of bipartite grouping and mass-weighted merging lets the task loss train these embeddings. DTEM supports both modular training with a frozen backbone and end-to-end training; in the latter, it alternates soft-merging updates for the embeddings with hard-merging updates for the backbone, then uses hard matching at inference. MergeNet also learns merge decisions through soft assignments, but restricts transport to the original patch grid and gathers a fixed number of carriers within the same training forward pass.

Spatial structure and recovery. DSM and ALGM combine local and global merging (Heo et al., 2024; Norouzi et al., 2024); GTP-ViT propagates tokens over spatial and semantic edges (Xu et al., 2024). Recent reducers constrain matching along rows/columns or after space-filling reordering (Gong & Lis, 2025; Li & Xu, 2025). MergeNet instead constrains eligible routes by Euclidean radius on the original grid. LookupViT uses cross-attention between compressed and high-resolution streams (Koner et al., 2024); MergeNet lets its selected carriers read the full-grid local features before latent processing.

3 METHOD

3.1 OVERVIEW

MergeNet connects differentiable routing to a physically shorter latent sequence within one forward pass. Given N image patches and a class token, a local encoder produces full-grid features $H^{(1)}, \dots, H^{(L)}$. Starting from $X^{(0)} = H^{(L)}$ and unit patch masses, L soft routing steps redistribute information while retaining every patch slot. A hard gather then selects K patch carriers. The carriers read the pre-transport features through cross-attention, and a latent encoder followed by a class head produces the prediction. Figure 1 sketches the pipeline.




Figure 1: Overview of the local encoder, spatial routing, fixed-budget selection, recovery read, and latent encoder in MergeNet.

3.2 SPATIAL MASS ROUTING

Spatial routing restricts where patch information may move without fixing the routing weights. Patch i keeps its coordinate u_i on the original two-dimensional grid. At routing step ℓ , patches are partitioned into disjoint donor and receiver sets A_ℓ and B_ℓ ; the class token participates in neither set. The eligible receivers for donor i are

$$\mathcal{E}_i^{(\ell)}(R) = \{j \in B_\ell : \|u_i - u_j\|_2 \leq R\}. \quad (1)$$

Coordinates remain tied to original slots as features move. The radius bounds each transport step; information may pass through multiple steps.

The routing weights combine learned similarity with a soft donor decision. A learned projection of $H^{(\ell)}$ yields normalized metrics $z_i^{(\ell)}$. For eligible pairs, we compute

$$s_{ij}^{(\ell)} = \langle z_i^{(\ell)}, z_j^{(\ell)} \rangle / \tau, \quad q_{ij}^{(\ell)} = \frac{\exp s_{ij}^{(\ell)}}{\sum_{r \in \mathcal{E}_i^{(\ell)}(R)} \exp s_{ir}^{(\ell)}}. \quad (2)$$

The row energy $e_i^{(\ell)} = \sum_{j \in \mathcal{E}_i^{(\ell)}(R)} \exp s_{ij}^{(\ell)}$ enters a threshold-based smooth top- k operator (Su, 2024). Given the step budget k_ℓ , it produces a continuous donor weight $g_i^{(\ell)}$; the transport weight is $a_{ij}^{(\ell)} = g_i^{(\ell)} q_{ij}^{(\ell)}$ on eligible edges and zero elsewhere; a donor with no eligible receiver has zero assignment. The smooth weights provide a gradient path from the task loss to the routing metrics. Appendix A gives the relaxation used here and for final carrier selection.

Soft transport updates mass and receiver features without removing the donor slot. With patch feature X_i and mass m_i , one step gives

$$\begin{aligned} m'_i &= m_i \left(1 - \sum_j a_{ij} \right), & X'_i &= X_i, \\ m'_j &= m_j + \sum_i a_{ij} m_i, & X'_j &= \frac{m_j X_j + \sum_i a_{ij} m_i X_i}{m'_j}. \end{aligned} \quad (3)$$

The sums cover eligible edges in the current donor–receiver partition. The implementation stabilizes residual fractions and denominators, while all patch slots remain present for the next routing step.

3.3 PHYSICAL BOTTLENECK AND RECOVERY

A hard gather turns the soft routing state into the shorter sequence used by later blocks. For compression factor λ , the patch-carrier budget is

$$K = N - \lfloor N(\lambda - 1)/\lambda \rfloor. \quad (4)$$

Let $\mathcal{I} = \text{TopK}(m_{\text{patch}}^{(L)}, K)$ contain the indices of the K largest patch masses, ordered before gathering. After normalizing the routed features, $\bar{X} = \text{Norm}(X^{(L)})$, we form the carrier sequence $C = [\bar{X}_{\text{cls}}; \bar{X}_{\mathcal{I}}]$. Thus C contains exactly $K + 1$ tokens, including the protected class token, in both training and inference.

The gather keeps its indices discrete while passing gradients to soft selection scores. A second smooth top- K operator applied to the final patch masses gives weights h_i . Each selected feature is multiplied by $1 + h_i - \text{stopgrad}(h_i)$, which equals one in the forward pass. This connects the selected features to the continuous scores without differentiating the index choice.

Recovery lets the carriers read information still available on the full grid. Using $H^{(L)}$ as the pre-transport key and value sequence, define

$$Z = C + \text{CrossAttn}(C, H^{(L)}, H^{(L)}). \quad (5)$$

The latent encoder processes Z , and its class-token output enters the classification head. Global latent self-attention therefore runs on $K + 1$ tokens, while recovery still reads the full-grid features.

3.4 MASS-AWARE LATENT ATTENTION

The selected carriers represent different amounts of transported mass. Following proportional attention in ToMe (Bolya et al., 2023), latent self-attention adds the log of each carrier mass to its key logit. For per-head query, key, and value matrices Q_a, K_a, V_a , head width d_h , and selected masses m_j , this is

$$\text{Attn}(Q_a, K_a, V_a; m) = \text{softmax} \left(\frac{Q_a K_a^\top}{\sqrt{d_h}} + \mathbf{1} \log(m)^\top \right) V_a. \quad (6)$$

The class token retains unit mass. This bias accounts for mass in attention; its contribution to classification accuracy requires a separate ablation.

The per-key bias can be encoded as one extra query/key dimension. Set $\hat{Q}_a = [Q_a, \sqrt{d_h} \mathbf{1}]$ and $\hat{K}_a = [K_a, \log m]$. With the original scale $1/\sqrt{d_h}$,

$$\frac{\hat{Q}_a \hat{K}_a^\top}{\sqrt{d_h}} = \frac{Q_a K_a^\top}{\sqrt{d_h}} + \mathbf{1} \log(m)^\top. \quad (7)$$

The released FlashAttention-compatible operator applies this identity after padding the augmented head dimension.

4 EXPERIMENTS

We evaluate MergeNet on ImageNet-1K classification, then compare routing variants to study the spatial prior. We also examine how accuracy changes with the carrier budget and image resolution, and distinguish theoretical savings from measured inference time.

4.1 EXPERIMENTAL SETUP

Data and training. Our main experiments train on ImageNet-1K and report EMA top-1 accuracy on all 50,000 validation images. At 224×224 resolution, the DeiT-S/8 backbone (Touvron et al., 2021) starts from 784 patch tokens. Dense DeiT-S/8 and MergeNet train from scratch for 300 epochs with AdamW, an effective global batch of 1,024, cosine learning-rate decay, five warmup epochs, and the same augmentation family. The selected MergeNet run uses a learning rate of 7.5×10^{-4} ; dense DeiT uses 5×10^{-4} . The spatial-routing controls use the latter rate for 150 epochs. Appendix A gives the full recipes and run inventory.

Comparisons. MergeNet gathers 392 patch carriers after six full-grid local blocks. ToMe and PiToMe instead compress the same trained dense checkpoint at inference time. We set all three methods to 392 final patches, excluding CLS, while keeping their different compression schedules explicit. The patch-8 DTEM adaptation follows the common recipe (Table 1). For the resolution study, 384-pixel models fine-tune the 224-pixel EMA weights for 30 epochs, and 512-pixel models receive a further 15 epochs.

Evaluation. The compressed rows use full-validation evaluations of selected EMA checkpoints; the dense score is the best EMA training-summary endpoint. In particular, MergeNet scores 81.360% at 392 patches in the re-evaluation, compared with 81.368% for its best training-summary checkpoint. The reported values are checked against training summaries and supplied evaluation records. The trained-checkpoint runtime comparison uses a common GPU, precision, and batch protocol, with each method’s accuracy-evaluation attention setting; existing synthetic benchmarks are reported separately in Section 4.4.

4.2 MAIN RESULTS

At the 392-patch operating point, MergeNet reaches 81.360% ImageNet top-1 accuracy (Table 1). ToMe and PiToMe, applied to the dense checkpoint after training, reach 81.932% and 81.742%; uncompressed DeiT-S/8 reaches 82.246%. Thus the proposed model learns its routing policy while processing a physically shorter latent sequence, although the available ToMe configuration retains more classification accuracy at the same final patch count. The matched DTEM-p8 adaptation reaches [VALUE]% top-1.

Table 1: ImageNet-1K top-1 accuracy at 224 pixels. Compressed rows target 392 final patches, excluding CLS.

Method	Checkpoint / path	Final patches	Top-1 (%)
DeiT-S/8	trained dense	784	82.246
DeiT-S/8 + ToMe	post-training	392	81.932
DeiT-S/8 + PiToMe	post-training	392	81.742
MergeNet ($R = 3$)	trained bottleneck	392	81.360
DTEM-p8	common-recipe adaptation	392	[VALUE]

The budget sweep shows how these models behave away from the default operating point (Figure 2). Using the same MergeNet checkpoint, increasing the final budget from 196 to 523 patches raises top-1 from 80.228% to 81.380%. ToMe remains ahead at each of the five 224-pixel budgets and at the five 384-pixel fine-tuned budgets shown in the figure. The complete sweep contains 20 matched final-budget settings, all with the same ordering (Appendix B). These are evaluations of fixed checkpoints; they describe accuracy as the inference budget changes rather than the outcome of retraining at each budget. The 384-pixel sweep uses a separately fine-tuned MergeNet checkpoint with $R = 3 \times 384/224 \approx 5.143$; the 384-pixel result in Table 2 uses the $R = 3$ fine-tune.

Fine-tuning at higher resolution reduces the gap to dense DeiT at 384 pixels, but the gap grows slightly at 512 pixels (Table 2). The best EMA difference changes from -0.878 points at 224 pixels to -0.394 at 384 pixels and -0.450 at 512 pixels. The last MergeNet model descends from the radius-scaled 384-pixel branch, so this sequence describes the selected configurations rather than one unchanged-radius trajectory.




Figure 2: ImageNet top-1 across final patch budgets at 224 pixels (left) and after 384-pixel fine-tuning (right). The latter MergeNet checkpoint uses $R = 5.143$. Lines connect evaluations of fixed checkpoints.

Table 2: Best EMA ImageNet-1K top-1 (%) across resolutions. Gap is MergeNet minus dense in percentage points.

Pixels	Schedule	DeiT-S/8 (%)	MergeNet (%)	Gap (pp)
224	300 scratch	82.246	81.368	-0.878
384	300+30	82.976	82.582	-0.394
512	300+30+15	83.068	82.618	-0.450

4.3 ABLATION STUDIES

Among the tested MergeNet routing variants, the original-grid radius-3 configuration scores above the global and flattened-window configurations. With the same 150-epoch ImageNet recipe and learning rate, radius 3 reaches 78.888% top-1, compared with 77.218% for global routing and 78.290% for a flattened width-8 window (Table 3). Radius 2 is slightly higher at 78.942%; these single-seed results identify a useful routing configuration, not an optimal radius. The global arm retains MergeNet’s architecture and is an internal variant rather than a DTEM reproduction. The variants also use their respective scoring implementations, detailed in Appendix A.

Table 3: ImageNet routing variants at 224 pixels. All rows use the same 150-epoch recipe and learning rate 5×10^{-4} .

Routing / probe	LR (10^{-4})	Best (%)	Final (%)
Global	5.0	77.218	77.216
Flat window 8	5.0	78.290	78.278
Spatial $R = 2$	5.0	78.942	78.942
Spatial $R = 3$	5.0	78.888	78.888

Radius also changes the number of eligible receivers. We therefore add a CIFAR-100 control that permutes radius-3 receivers while preserving donor candidate counts and receiver degrees (Appendix A).

Across three seeds on resized CIFAR-100, radius 3 has the highest epoch-199 EMA top-1 mean among the four routing policies (Table 4). Its paired gain over the degree-matched control is 1.42 ± 0.30 points, with all three seed-wise differences positive. This supports the spatial prior within the tested setting; the replicated result does not estimate ImageNet seed variance.

Table 4: CIFAR-100 epoch-199 EMA top-1 over seeds 42–44. Δ is the paired $R = 3$ minus control difference (pp).

Routing	Top-1 (%)	$R = 3$ minus control (pp)
Global	68.48 $\pm$ 0.10	+2.11 $\pm$ 0.87
Flat window 8	69.29 $\pm$ 0.54	+1.30 $\pm$ 0.47
Degree-matched	69.17 $\pm$ 0.62	+1.42 $\pm$ 0.30
Spatial $R = 3$	70.59 $\pm$ 0.88	—

The component study tests whether the trained classifier depends on mass-aware attention and full-grid recovery. With the same checkpoint, carrier selection, and 50,000 validation images, disabling the $\log m$ key bias changes top-1 by [VALUE] percentage points; bypassing residual recovery changes it by [VALUE] points, and disabling both changes it by [VALUE] points (Appendix D). These inference-only interventions measure dependence of the trained model, rather than the effect of training without each component.

4.4 EFFICIENCY

At 224 pixels, MergeNet passes 392 rather than 784 patches to the latent stage and requires an estimated 14.13 GMAC per image, versus 22.41 GMAC for dense DeiT-S/8. The estimate excludes indexing and data movement.

Table 5 and Figure 3 summarize the random-weight A100 measurements. Appendix C gives the protocol and compiled results.

The common trained-checkpoint evaluation gives MergeNet [VALUE] ms batch-1 latency, [VALUE] images/s batch-64 throughput, and [VALUE] GiB peak allocated memory (Appendix D).

Table 5: Synthetic A100 inference at 224 pixels with fp16 autocast and random weights. B1/B64 denote batch sizes 1 and 64; B64 memory is peak allocated and reserved GiB. The MergeNet B1 range spans two independent run medians.

Method	Patches	B1 (ms)	B64 (ms)	B64 (im/s)	Alloc.	Reserv.
Dense DeiT-S/8	784	4.38	30.47	2100.5	0.744	0.957
ToMe	392	12.03	34.72	1843.5	0.724	1.004
PiToMe	392	14.17	70.28	910.7	1.774	3.104
MergeNet ($R = 3$)	392	18.27–21.00	36.30	1763.3	0.776	0.836




Figure 3: Batch-64 eager inference on the A100 with random weights. Bars show median latency and peak allocated memory from the same measurements as Table 5; lower is better.

4.5 VISUALIZATION

The spatial prior also makes routing behavior inspectable. Figure 4 shows donor-to-receiver edges and final carrier locations for fixed ImageNet examples, while Figure 5 summarizes selection frequency and routing distance over 1,000 images. The passive-trace check gives [logit, carrier-index, and mass parity] (Appendix E).

5 DISCUSSION

The results separate the routing mechanism from overall model accuracy. The exact gather shortens the latent sequence, and the original-grid variant scores above the tested global and flattened-window variants on ImageNet and three-seed CIFAR-100. At equal final patch counts, the complete classifier trails ToMe on ImageNet. The routing-variant and cross-method comparisons therefore answer different questions.

Deployment cost and component attribution require separate measurements from the carrier count. Full-grid processing and routing overhead affect runtime; the trained-checkpoint comparison gives [\[latency and throughput interpretation\]](#). ImageNet geometry uses one seed, while the replicated locality test uses resized CIFAR-100. The fixed-checkpoint component tests give [\[mass-bias and recovery interpretation\]](#). The routing visualization describes transport without implying semantic boundary preservation.

6 CONCLUSION

MergeNet combines differentiable spatial transport with an exact-budget gather so that learned routing feeds a shorter latent sequence. The radius-3 variant scores above the tested global and flattened variants; a three-seed comparison also favors it over degree-matched routing. The current ImageNet result is 81.360% top-1 at 392 final patches, with a clear accuracy gap to ToMe at that budget. Under a common inference protocol, the trained-checkpoint accuracy–efficiency tradeoff is [\[summary relative to baselines\]](#). Reducing full-grid and routing overhead remains a direction for further work.

REFERENCES

- Daniel Bolya, Cheng-Yang Fu, Xiaoliang Dai, Peizhao Zhang, Christoph Feichtenhofer, and Judy Hoffman. Token Merging: Your ViT but Faster. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.09461v3>.
- Mengzhao Chen, Wenqi Shao, Peng Xu, Mingbao Lin, Kaipeng Zhang, Fei Chao, Rongrong Ji, Yu Qiao, and Ping Luo. DiffRate: Differentiable Compression Rate for Efficient Vision Transformers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2023. URL <https://arxiv.org/abs/2305.17997v1>.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. In *Advances in Neural Information Processing Systems*, 2022. URL <https://proceedings.neurips.cc/paper/2022/hash/67d57c32e20fd0a7a302cb81d36e40d5-Abstract-Conference.html>.
- Wenyi Gong and Mieszko Lis. CubistMerge: Spatial-Preserving Token Merging for Diverse ViT Backbones, 2025. URL <https://arxiv.org/html/2509.21764v2>. arXiv version 2, March 2, 2026; first posted September 26, 2025.
- Jung Hwan Heo, Seyedarmin Azizi, Arash Fayyazi, and Massoud Pedram. Training-Free Acceleration of ViTs with Delayed Spatial Merging, 2024. URL <https://arxiv.org/abs/2303.02331v2>. ICML 2024 ES-FoMo Workshop; arXiv version 2, July 1, 2024.
- Rajat Koner, Gagan Jain, Prateek Jain, Volker Tresp, and Sujoy Paul. LookupViT: Compressing Visual Information to a Limited Number of Tokens. In *European Conference on Computer Vision*, 2024. URL <https://arxiv.org/html/2407.12753v1>.
- Dong Hoon Lee and Seunghoon Hong. Learning to Merge Tokens via Decoupled Embedding for Efficient Vision Transformers. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/html/2412.10569v1>.
- Sanghyeok Lee, Joonmyung Choi, and Hyunwoo J. Kim. Multi-criteria Token Fusion with One-step-ahead Attention for Efficient Vision Transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/papers/Lee_Multi-criteria_Token_Fusion_with_One-step-ahead_Attention_for_Efficient_Vision_Transformers_CVPR_2024_paper.pdf.

- Yunge Li and Lanyu Xu. Neighbor-Aware Token Reduction via Hilbert Curve for Vision Transformers, 2025. URL <https://arxiv.org/html/2512.22760v1>. arXiv version 1, December 28, 2025.
- Narges Norouzi, Svetlana Orlova, Daan de Geus, and Gijs Dubbelman. ALGM: Adaptive Local-then-Global Token Merging for Efficient Semantic Segmentation with Plain Vision Transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. URL https://openaccess.thecvf.com/content/CVPR2024/papers/Norouzi_ALGM_Adaptive_Local-then-Global-Token-Merging_for_Efficient_Semantic_Segmentation_with_CVPR_2024_paper.pdf.
- Yongming Rao, Wenliang Zhao, Benlin Liu, Jiwen Lu, Jie Zhou, and Cho-Jui Hsieh. DynamicViT: Efficient Vision Transformers with Dynamic Token Sparsification. In *Advances in Neural Information Processing Systems*, 2021. URL https://papers.nips.cc/paper_files/paper/2021/file/747d3443e319a22747fbb873e8b2f9f2-Paper.pdf.
- Jianlin Su. Softmax Backpropagation: Searching for a Smooth Approximation to Top- k , September 2024. URL <https://spaces.ac.cn/archives/10373>. Scientific Spaces blog; accessed September 22, 2026.
- Hugo Touvron, Matthieu Cord, Matthijs Douze, Francisco Massa, Alexandre Sablayrolles, and Herve Jegou. Training data-efficient image transformers & distillation through attention. In *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pp. 10347–10357. PMLR, 2021. URL <https://proceedings.mlr.press/v139/touvron21a.html>.
- Hoai-Chau Tran, Duy M. H. Nguyen, Duy M. Nguyen, Trung-Tin Nguyen, Ngan Le, Pengtao Xie, Daniel Sonntag, James Y. Zou, Binh T. Nguyen, and Mathias Niepert. Accelerating Transformers with Spectrum-Preserving Token Merging. In *Advances in Neural Information Processing Systems*, 2024. URL <https://arxiv.org/abs/2405.16148v2>.
- Weiying Xie, Xiaoyu Chen, Xin Zhang, Chenhe Hao, Jitao Ma, Yunsong Li, and Leyuan Fang. Saliency-Driven Token Merging for Vision Transformers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026. URL https://openaccess.thecvf.com/content/CVPR2026/html/Xie_Saliency-Driven-Token-Merging_for_Vision_Transformers_CVPR_2026_paper.html.
- Xuwei Xu, Sen Wang, Yudong Chen, Yanping Zheng, Zhewei Wei, and Jiajun Liu. GTP-ViT: Efficient Vision Transformers via Graph-Based Token Propagation. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2024. URL https://openaccess.thecvf.com/content/WACV2024/papers/Xu_GTP-ViT_Efficient_Vision_Transformers_via_Graph-Based-Token_Propagation_WACV_2024_paper.pdf.
- Hongxu Yin, Arash Vahdat, Jose M. Alvarez, Arun Mallya, Jan Kautz, and Pavlo Molchanov. A-ViT: Adaptive Tokens for Efficient Vision Transformer. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022. URL https://openaccess.thecvf.com/content/CVPR2022/papers/Yin_A-ViT_Adaptive_Tokens_for_Efficient_Vision_Transformer_CVPR_2022_paper.pdf.

A IMPLEMENTATION DETAILS

Table 6 gives the training settings for the main ImageNet experiments. All runs use an effective global batch of 1,024. The 150-epoch routing variants share a learning rate of 5×10^{-4} ; the selected 300-epoch MergeNet run uses 7.5×10^{-4} . EMA top-1 is measured on the 50,000-image validation set.

Table 6: Training settings for the ImageNet-1K experiments.

Setting	224-pixel scratch	384/512-pixel fine-tune
Epochs	150 or 300	30 at 384; 15 at 512
Effective global batch	1024	1024
Optimizer	AdamW	AdamW
Learning rate	5×10^{-4} or 7.5×10^{-4}	10^{-5}
Minimum learning rate	10^{-5}	10^{-6}
Weight decay	0.05	10^{-8}
Warmup epochs	5	1
Warmup learning rate	10^{-6}	10^{-6}
EMA decay	0.99996	0.9996
Mixup / CutMix	0.8 / 1.0	0 / 0
Label smoothing	0.1	0.1
Random erasing	0.25	0.25
RandAugment	rand-m9-mstd0.5-inc1	same
Gradient clip norm	1.0	1.0
Validation crop percentage	0.9	1.0
MergeNet auxiliary-logit weight	ramp to 0.05	0.05
Teacher / feature / routing distillation	disabled	disabled

Classifier. The patch-8 model has width 384, six heads, six local blocks, and six latent blocks. At 224 pixels, the local encoder processes 784 patch slots. Local attention spans 16 sequence positions on either side of a patch query; the CLS query attends globally. Spatial routing uses radius 3, a $384 \rightarrow 64$ metric projection, similarity temperature 1, and metric input-gradient scale 0.1. At compression factor 2, the six soft donor budgets are (65, 65, 65, 65, 65, 67); the hard gather retains 392 patches. The latent encoder does not merge further. Stochastic depth is 0.1, with separate local and latent ramps.

Routing and selection. Training partitions donors and receivers randomly for each image; evaluation alternates deterministic odd/even partitions across routing steps. The final gather indices are detached, and selected carriers are ordered by a detached, mass-weighted center of their original flattened indices. Spatial scoring uses sparse neighbor lists; the global and flattened-window controls use their respective historical dense-scoring paths.

Smooth top- k . For donor energies e_i and step budget k_ℓ , let $v_i = 30e_i/(\max e - \min e)$. A threshold θ is determined from the sorted scores using prefix and suffix exponential sums. The continuous donor weight is

$$g_i = \begin{cases} 1 - \frac{1}{2} \exp(\theta - v_i), & v_i > \theta, \\ \frac{1}{2} \exp(v_i - \theta), & v_i \leq \theta. \end{cases} \quad (8)$$

The threshold makes the weights sum to k_ℓ up to numerical precision. Donors without eligible receivers are assigned zero transport weight. Applying the same relaxation to the final patch masses supplies the straight-through carrier weights; the hard top- K indices still determine which slots are gathered.

At training time, an auxiliary full-grid readout contributes up to 0.05 to the latent CLS logits. Its weight begins ramping at epoch 10 for the 150-epoch runs and epoch 20 for the 300-epoch run, reaching 0.05 over 10 and 20 epochs, respectively. It is disabled at evaluation.

CIFAR-100 controls. The three-seed routing comparison uses 12 independent 200-epoch runs: global, flattened-window, radius-3, and degree-matched routing for seeds 42–44. Native 32×32 images are resized to 224 pixels. The primary endpoint is epoch-199 EMA top-1 on the 10,000-image test set. Each training and evaluation batch contains 200 images. The degree-matched arm permutes receiver columns of the realized radius-3 support using fixed receiver priorities; it preserves each donor’s candidate count and the receiver-degree multiset. Table 4 reports means and sample standard deviations across seeds, with paired same-seed differences.

B ACCURACY BUDGETS

Table 7 gives the values behind Figure 2. The MergeNet checkpoint uses radius 3 at 224 pixels and radius 5.143 after 384-pixel fine-tuning. Each row matches final patch counts, excluding CLS, but not the layerwise sequence lengths: ToMe and PiToMe merge across blocks, whereas MergeNet gathers once between stages. All points evaluate fixed checkpoints at different inference budgets.

Table 7: ImageNet top-1 (%) at matched final patch counts. Δ is ToMe minus MergeNet in percentage points.

Checkpoint / pixels	Patches	MergeNet	ToMe	PiToMe	Δ ToMe
224 trained / 224	523	81.380	82.038	81.876	+0.658
224 trained / 224	392	81.360	81.932	81.742	+0.572
224 trained / 224	314	81.074	81.814	81.560	+0.740
224 trained / 224	262	80.850	81.830	81.500	+0.980
224 trained / 224	196	80.228	81.752	81.206	+1.524
384 tuned / 384	1537	82.500	82.868	82.658	+0.368
384 tuned / 384	1152	82.562	82.866	82.534	+0.304
384 tuned / 384	923	82.230	82.800	82.452	+0.570
384 tuned / 384	770	82.002	82.792	82.334	+0.790
384 tuned / 384	576	81.300	82.742	82.296	+1.442

C SYNTHETIC EFFICIENCY

We measure inference on one idle A100-SXM4-80GB with randomly initialized models, device-resident 224×224 inputs, and fp16 autocast. Each method and batch size runs in a fresh process with 50 warmups followed by five rounds of 200 CUDA-event samples. The reported latency is the median of all 1,000 samples; timing excludes model setup, data loading, and diagnostic hooks. Dense DeiT-S/8 has 784 patch tokens, while each compressed method reaches 392 final patches, excluding CLS. ToMe and PiToMe both enable proportional attention.

Table 5 reports the eager measurements. The five-round median coefficient of variation is below 0.3% for every batch-64 cell. Batch-1 measurements are more sensitive to host scheduling: two independent MergeNet runs yield 18.27 and 21.00 ms, despite each passing the 3% within-run variation threshold.

As a separate implementation comparison, we compile ToMe’s transformer block forwards. For MergeNet, we compile local and latent transformer blocks, cross-attention, and selection and merge operations, while leaving the enclosing local-encoder forward eager. At batch 64, their latencies are 27.22 and 27.24 ms, respectively; peak allocated memory is 0.614 and 0.627 GiB. At batch 1, ToMe takes 9.24 ms and MergeNet takes 11.22 ms. The compiled and eager measurements are separate operating points. All A100 measurements use random weights and cannot form an accuracy–latency comparison with trained checkpoints. The current ToMe and PiToMe adapters also do not use the fixed upstream commits required for the final common-protocol evaluation (Appendix D).

D CHECKPOINT EVALUATION

The common inference benchmark uses the same GPU, precision, and measurement procedure for all methods. Each method retains its accuracy-evaluation attention setting, and achieved token counts are recorded. The benchmark reports full-validation accuracy separately from batch-1

latency, batch-64 throughput, and peak allocated memory. Table 8 gives the trained-checkpoint measurements; the random-weight A100 timings in Appendix C are reported separately.

Table 8: Common-protocol trained-checkpoint inference at 224 pixels. B1/B64 denote batch sizes 1 and 64.

Method	Patches	B1 (ms)	B64 (im/s)	B64 alloc. (GiB)
Dense DeiT-S/8	784	[VALUE]	[VALUE]	[VALUE]
DTEM-p8	392	[VALUE]	[VALUE]	[VALUE]
DeiT-S/8 + ToMe	392	[VALUE]	[VALUE]	[VALUE]
DeiT-S/8 + PiToMe	392	[VALUE]	[VALUE]	[VALUE]
MergeNet ($R = 3$)	392	[VALUE]	[VALUE]	[VALUE]

The component test uses the same trained MergeNet EMA checkpoint for all four rows of Table 9. The mass-bias intervention sets only the latent log-mass key bias to zero; the recovery intervention bypasses only residual cross-attention. The native score is checked against the common inference evaluation. Paired intervals compare per-image correctness on the same 50,000 validation examples and do not represent training-seed uncertainty.

Table 9: Inference-only component test at 224 pixels. Δ and its paired 95% interval are relative to native top-1.

Inference setting	Top-1 (%)	Δ top-1 (pp), 95% CI
Native (mass + recovery)	[VALUE]	—
Mass bias off	[VALUE]	[VALUE]
Recovery off	[VALUE]	[VALUE]
Both off	[VALUE]	[VALUE]

E ROUTING VISUALIZATION

The visual examples are selected before inspecting the router. We choose one ImageNet validation image per class using a fixed SHA-256 ordering of image paths, giving 1,000 images in total. From this set, the main figure shows the first correctly classified and the first misclassified image under the same ordering. Passive tracing is checked against untraced predictions, top- K indices, and token mass; the check gives [parity summary]. If no error occurs in the selected set, the selection rule uses two correct cases.

The main figure pairs actual step-3 donor-to-receiver transport edges with the final 392 carrier locations for each selected image. The population figure summarizes per-image routing distances over all six steps and the selection frequency of each original 28×28 patch location across the 1,000 images. Carrier markers denote selected slots, not semantic regions. Mean route distance is [VALUE] patches, and the selection-frequency pattern is [population interpretation].

[Figure: E6: actual routing edges and final carriers for fixed examples]

Figure 4: ImageNet routing examples: step-3 transport edges and final carrier locations from the same trained checkpoint.

[Figure: E6: routing distances and carrier selection frequency]

Figure 5: Population view for 1,000 fixed validation images. Distances are measured in original-grid patches; selection frequency uses the 392 final carrier slots.